

# Principles of Safe Autonomy

## ECE 484 SP'26

### Lecture 2

Professor: Sayan Mitra  
Jan 22, 2026

[https://safeautonomy-  
illinois.github.io/ece484-site/](https://safeautonomy-illinois.github.io/ece484-site/)

# Outline

Motivation

Administrivia

Introduction to Safety

- Models
- Requirements
- Proofs



# Automata or state machine models

An **automaton**  $A$  is defined by a triple  $\langle Q, Q_0, D \rangle$ , where

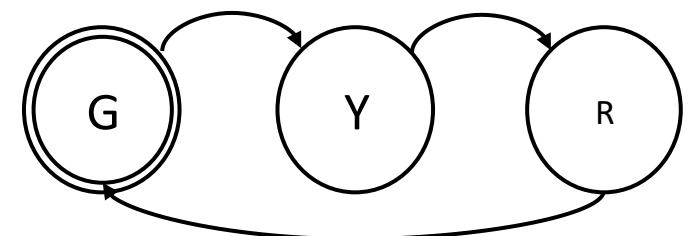
- ▶  $Q$  is a set of **states**
- ▶  $Q_0 \subseteq Q$  is a set of **initial states**
- ▶  $D \subseteq Q \times Q$  is a set of **transitions**

An **execution** of  $A$  is a finite or infinite sequence  $q_0, q_1, \dots$  such that  $q_0 \in Q_0$  and  $(q_i, q_{i+1}) \in D$

Example: Traffic light automaton

- ▶  $Q = \{G, Y, R\}$   $Q_0 = \{G\}$
- ▶  $D = \{(G, Y), (Y, R), (R, G)\}$

Execution of traffic light  $G, Y, R, G, Y, R \dots$  infinite even though finite state



# Requirements and Counter-examples

Requirements define what the system must and must not do

Example: “Car stays within speed limit”

Autonomous car: “Ego should not collide with lead car”

Collatz: “Every number eventually ends in the 4-2-1 cycle”

A **requirement** defines a set  $R$  of allowed executions

An execution  $\alpha$  that is not in the set  $R$  is a **counter-example**

$$R_{\text{eventually-1}} = \{\alpha \mid \exists k \alpha_k = 1\}$$

An automaton  $A$  **satisfies** a requirement  $R$  if *all* executions of  $A$  satisfies  $R$

Whether the Collatz automaton satisfies the requirement  $R_{\text{eventually-1}}$  for all initial conditions remains an open problem, although no counter-example has been found up to  $2^{70}$

This is an example of a **verification problem**

# Verification problem

**Verification problem:** Given an automaton  $A$  and a requirement  $R$ , check whether all executions of  $A$  satisfy  $R$  or find a counter-example

Testing or checking individual executions can help find counter-examples but cannot show that there is no counter-example

Verification can be hard because

- ▶  $|Q|$  is finite but large and testing may require visiting all the states (e.g., Collatz)
- ▶  $|Q|$  is small but the number of executions is very large
- ▶  $|Q|$  may be infinite and  $D$  may be nondeterministic --- typical for autonomous system

# Example: Automatic Emergency Braking (AEB)

Car must brake to maintain safe gap with lead vehicle/pedestrian

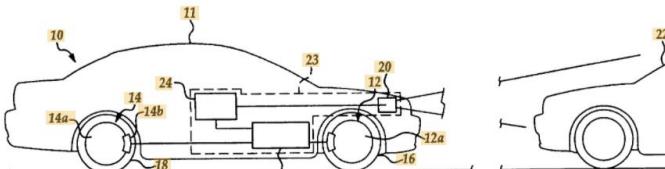


Figure 1

There is no standard for checking correctness of AEB systems

Future: Every code commit in github from an AEB engineer, **proves a theorem** establishing  $A$  satisfies  $R_{\text{gap}}$

[www.google.com](http://www.google.com) > patents

[US20110168504A1 - Emergency braking system - Google ...](https://patents.google.com/patent/US20110168504A1)

Jump to [Patent citations](#) (18) - [US4053026A](#) \* 1975-12-09 1977-10-11 Nissan Motor Co., Ltd. Logic circuit for an automatic braking system for a motor ...

[www.google.com](http://www.google.com) > patents

[US5170858A - Automatic braking apparatus with ultrasonic ...](https://patents.google.com/patent/US5170858A)

An automatic braking apparatus includes: an ultrasonic wave emitter provided in a ... Info: Patent citations (13); Cited by (7); Legal events; Similar documents; Priority and ... [US6523912B1](#) 2003-02-25 Autonomous emergency braking system.

[www.google.com](http://www.google.com) > patents

[DE102004030994A1 - Brake assistant for motor vehicles ...](https://patents.google.com/patent/DE102004030994A1)

B60T7/22 Brake-action initiating means for automatic initiation; for initiation not ... Info: Patent citations (3); Cited by (9); Legal events; Similar documents ... data from the environment sensor and then automatically initiates emergency braking.

[www.google.com.pg](http://www.google.com.pg) > patents

[Braking control system for vehicle - Google Patents](https://patents.google.com/patent/Braking%20control%20system%20for%20vehicle)

An automatic emergency braking system for a vehicle includes a forward viewing camera and a control. At least in part responsive to processing of captured ...

[www.automotiveworld.com](http://www.automotiveworld.com) > news-releases > toyota-ip... ▾

[Toyota IP Solutions and IUPUI issue first commercial license ...](https://www.automotiveworld.com/news-releases/toyota-ip/Toyota-IP-Solutions-and-IUPUI-issue-first-commercial-license)

Jul 22, 2020 - ... and validation of automotive automatic emergency braking (AEB) ... and Director of Patent Licensing for Toyota Motor North America. "We are ...

[insurancenewsnetwork.com](http://insurancenewsnetwork.com) > oarticle > patent-application-tit... ▾

[Patent Application Titled "Multiple-Stage Collision Avoidance" ...](https://www.insurancenewsnetwork.com/oarticle/patent-application-tit/Patent-Application-Titled-Multiple-Stage-Collision-Avoidance)

Apr 3, 2019 - No assignee for this patent application has been made. ... Automatic emergency braking systems will similarly, also, soon be required for tractor ...

# Automaton model of AEB

Automaton  $A = \langle Q, Q_0, D \rangle$

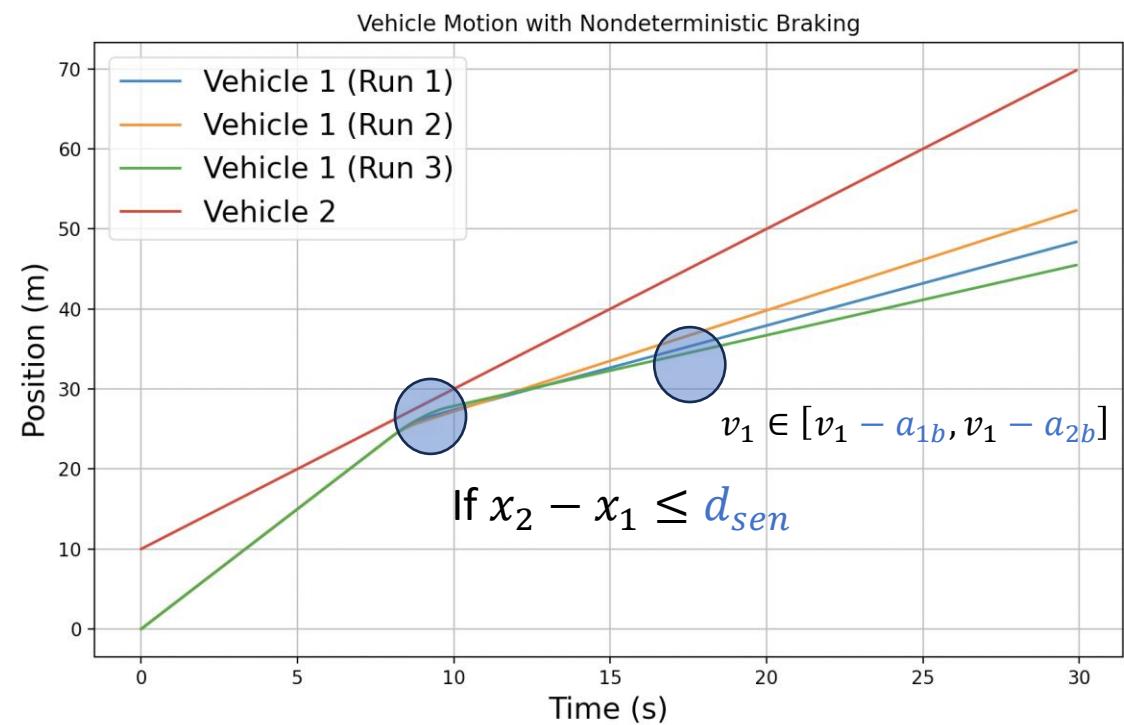
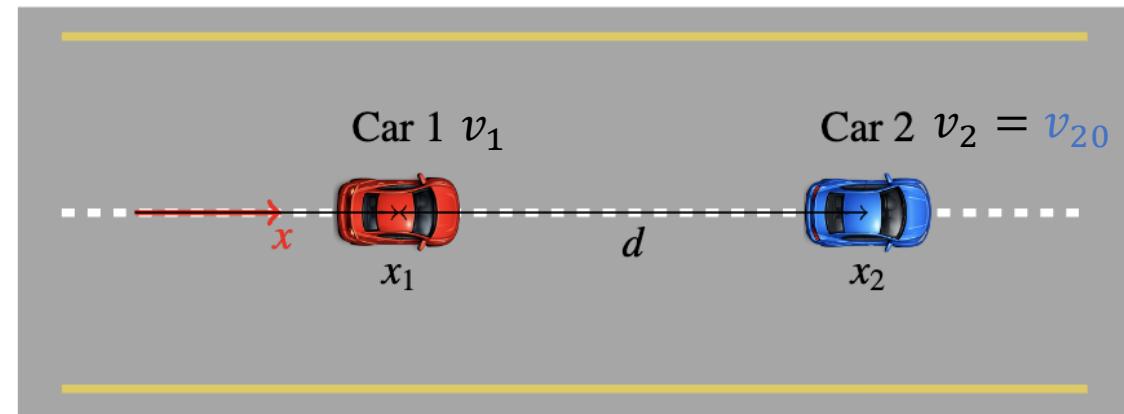
- $Q: [x_1, x_2, v_1] \in \mathbb{R}^3$
- $Q_0 = \{[x_1 = x_{10}, x_2 = x_{20}, v_1 = v_{10}]\}$
- $D \subseteq Q \times Q$  written as a program

If  $x_2 - x_1 \leq d_{sen}$

$$v_1 \in [v_1 - a_{1b}, v_1 - a_{2b}]$$

$$x_2 = x_2 + v_2$$

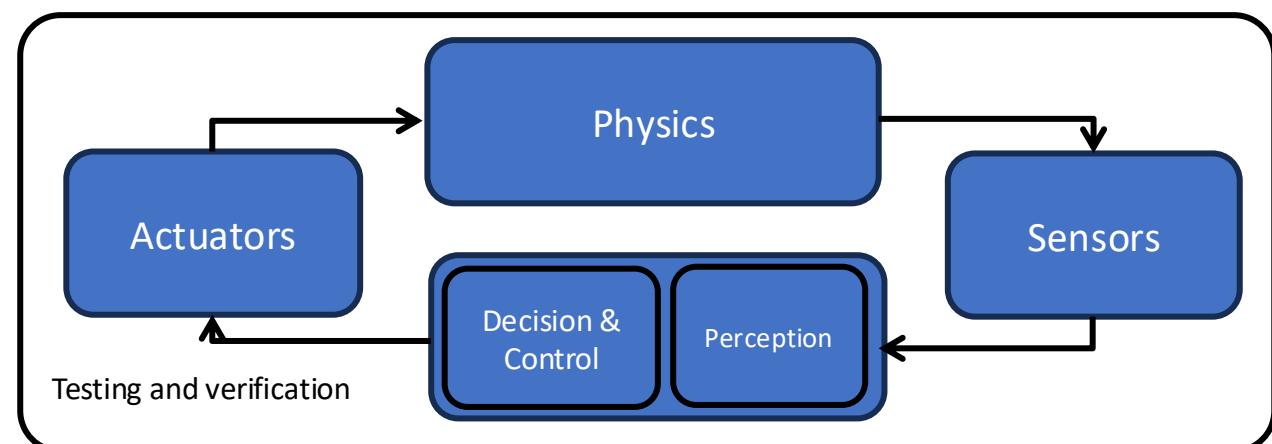
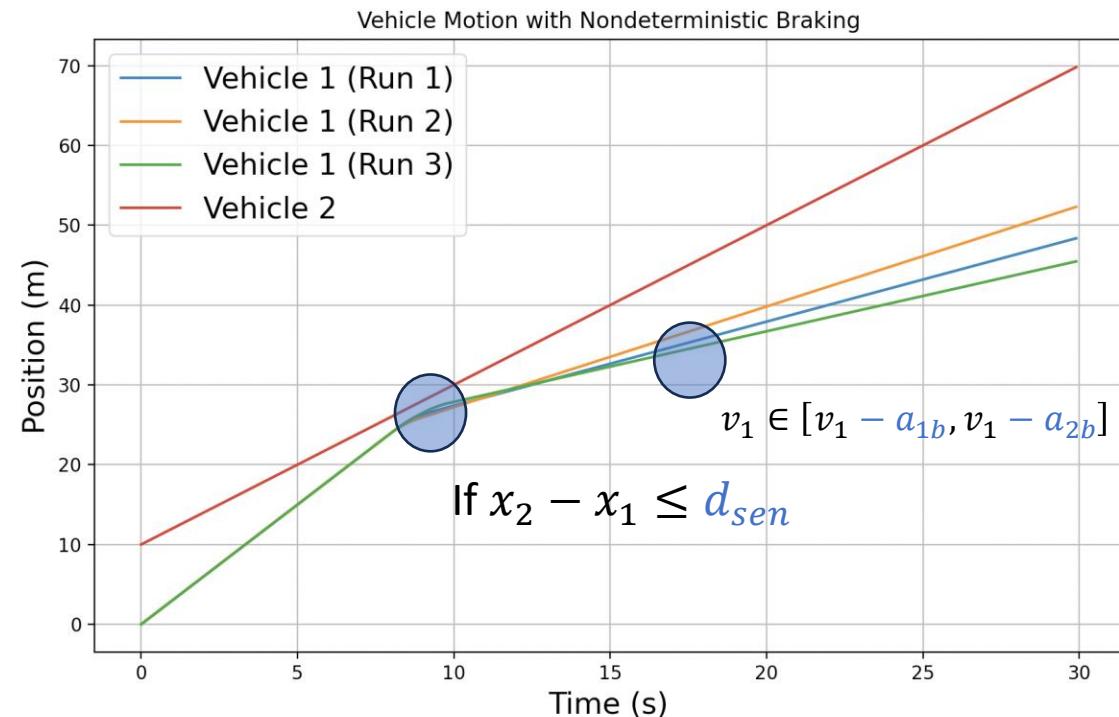
$$x_1 = x_1 + v_1$$



# Automaton model of AEB

Automaton  $A = \langle Q, Q_0, D \rangle$

- ▶  $Q: \mathbb{R}^3; q \in Q \quad q \cdot x_1, q \cdot x_2 \in \mathbb{R}$
- ▶  $Q_0 = \{q \mid [q \cdot x_1 = x_{10}, q \cdot x_2 = x_{20}, q \cdot v_1 = v_{10}]\}$
- ▶  $(q, q') \in D$  iff
  - if  $q \cdot x_2 - q \cdot x_1 \leq d_{sen}$
  - $q' \cdot v_1 \in [q \cdot v_1 - a_{1b}, q \cdot v_1 - a_{2b}]$
  - $q' \cdot x_2 = q \cdot x_2 + q \cdot v_2$
  - $q' \cdot x_1 = q \cdot x_1 + q \cdot v_1$



# What is missing in the AEB model?

If  $x_2 - x_1 \leq 2.0$

$v_1 \in [v_1 - a_{1b}, v_1 - a_{2b}]$

else  $v_1 = v_1$

$x_2 = x_2 + v_2$

$x_1 = x_1 + v_1$

- ▶ Acceleration, friction in dynamics
- ▶ Uncertainty in sensing
- ▶ Uncertainty in lead vehicle behavior
- ▶ Timing of execution of control loop

“All models are wrong, but some are useful.”

# Safety and liveness requirements

$$R_{gap} = \{\alpha \mid \forall i \alpha_i.x_2 > \alpha_i.x_1\}$$

non-zero gap  $U_{gap} = \{q \mid q.x_2 - q.x_1 \leq 0\}$

$$R_{sp-lim} = \{\alpha \mid \forall i \alpha_i.v_1 \leq 70\}$$

speed limit  $U_{sp-lim} = \{q \mid q.x_1 \geq 70\}$

$$R_{catch-up} = \{\alpha \mid \exists i 2 > \alpha_i.x_2 - \alpha_i.x_1 > 1\}$$

catch eventually

A **safety requirement** says that *every* state along *every* execution should stay in safe states

Equivalently, no execution of A ever reaches any unsafe states

$R_{gap}$  and  $R_{sp-lim}$  are safety requirements with  $U_{gap}$  and  $U_{sp-lim}$  as the unsafe sets

$R_{catch-up}$  is not a safety requirement; it is an example of a **liveness / progress requirement**

A liveness requirement says that along every execution eventually some good state is reached

# Outline

Motivation

Administrivia

Introduction to Safety

- Models
- Requirements
- Verification



# Safety verification: Finite State Automata

**Safety verification problem:** Given an automaton  $A$  and an unsafe set  $U$ , check whether there exists any execution  $\alpha$  of  $A$  that reaches  $U$

Counter-examples of safety are finite executions ending in  $U$

For finite automata, safety verification can be solved using depth first search from  $Q_0$

- ▶ Consider  $\langle Q, Q_0, D \rangle$  as a directed graph with  $\langle Q, D \rangle$
- ▶ DFS computes all paths or executions from  $Q_0$
- ▶ If none of these executions hit  $U$  there is no counter-example
- ▶ Absence of a counter-example **proves** that the automaton is safe

In practice, explicit enumeration of all paths may not scale to large graphs

# Safety verification and Reachability: Infinite State Spaces

A state  $q \in Q$  is **reachable** if there exists an execution  $\alpha$  such that  $\alpha_i = q$ .

$Reach_A(Q_0) \subseteq Q$  the set of reachable states of  $A$  from  $Q_0$

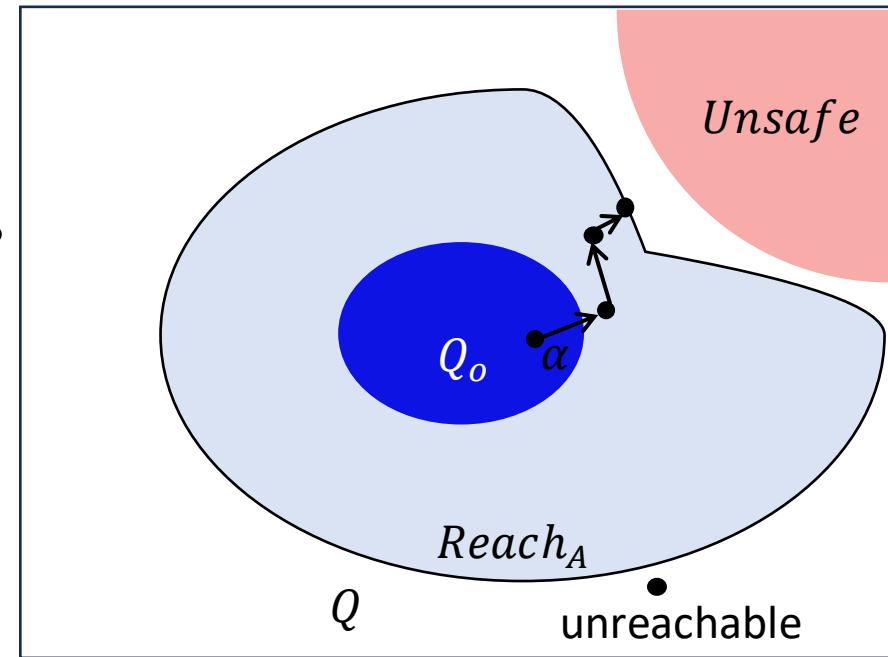
Safety verification problem is equivalent to as checking  $Reach_A \cap U = \emptyset$ ?

That is, if we can compute  $Reach_A$  then we can verify safety

Finite state systems DFS computes  $Reach_A(Q_0)$

For infinite state systems, we need:

- ▶ Representation of infinite sets of states
- ▶ Iteratively computing  $Reach_A$



# Computing reachable sets and over-approximations

Define  $\text{Post}(R) = \{q' \mid \exists q \in R, (q, q') \in D\}$  that gives all the states that can be reached in one step from the set of states  $R$

- ▶ For a deterministic system  $\text{Post}(\{q\}) = q'$  for  $(q, q') \in D$
- ▶ For finite  $R$ ,  $\text{Post}(R) = \bigcup_{q \in R} \text{Post}(\{q\})$
- ▶  $\text{Post}(Q_0) = \{q \mid \exists q_0 \in Q_0, (q_0, q) \in D\}$  states reachable in 1 step from  $Q_0$
- ▶  $\text{Post}(\text{Post}(Q_0)) = ?$

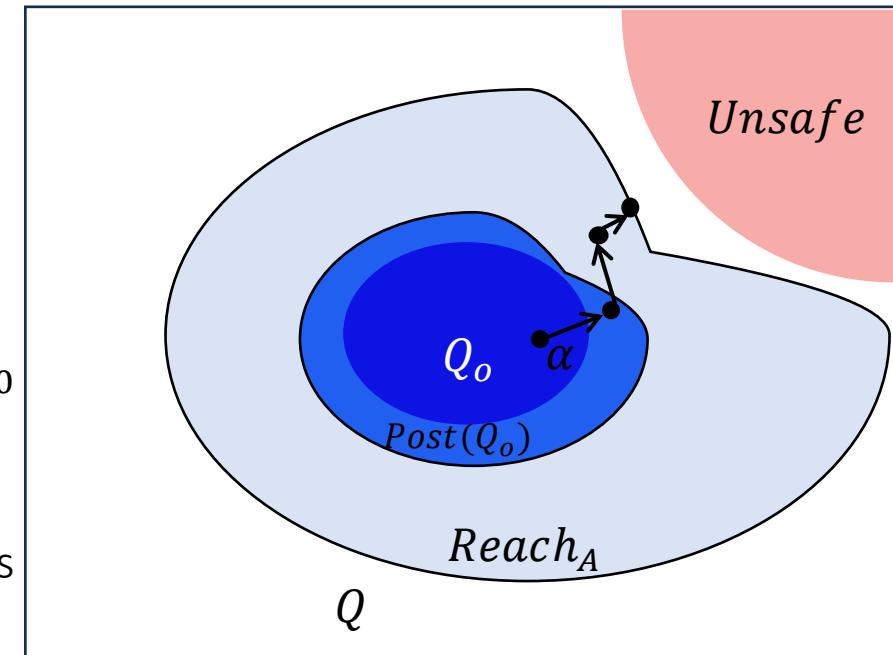
Infinite sets & nondeterministic  $\text{Post}(R)$  requires some representation of sets

Example:

- ▶  $Q = [x_1: \mathbb{R}] D: x_1 = x_1 + v_1$  then  $R = [a, b]$   $\text{Post}(R) = [a + v_1, b + v_1]$
- ▶  $Q = \mathbb{R}^4$  then  $R = [\mathbf{a}, \mathbf{b}]$  then  $\text{Post}(R)$  is a hyperrectangle

Generally, for nonconvex  $R$  nonlinear  $D$  exact  $\text{Post}(R)$  may be infeasible

We use over-approximation  $\overline{\text{Post}}(R)$  such that  $\text{Post}(R) \subseteq \overline{\text{Post}}(R)$



# Reachable sets and over-approximations

```
Reachability( $A = \langle Q, Q_0, D \rangle$ )
```

```
 $R_0 = Q_0$ 
```

```
 $R_1 = \emptyset$ 
```

```
 $i = 0$ 
```

```
do
```

```
   $R_{i+1} = \overline{Post}(R_i) \cup R_i$ 
```

```
   $i = i + 1$ 
```

```
Until  $R_i \neq R_{i-1}$ 
```

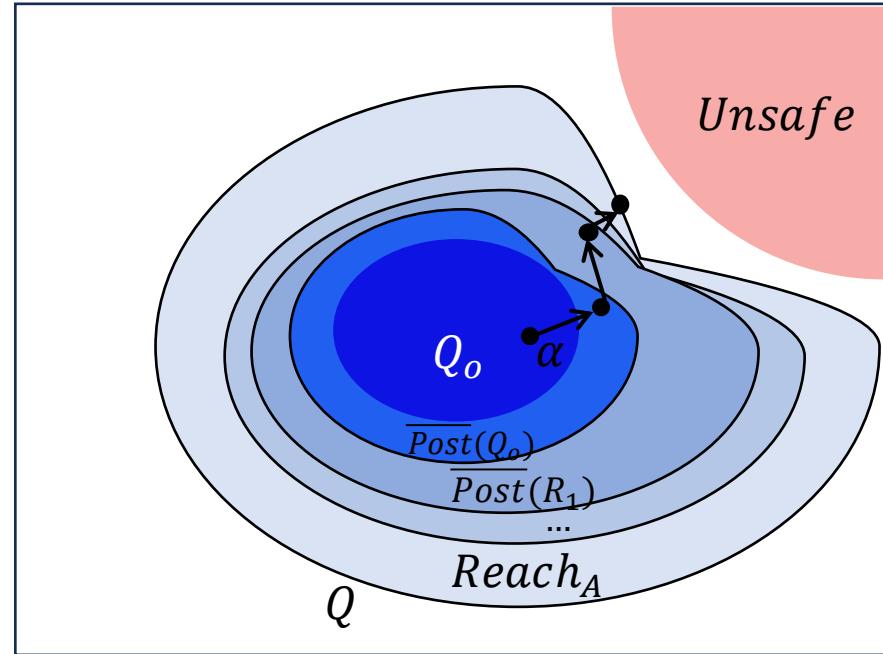
```
Return  $R_i$ 
```

Exercise. Show that  $\text{Post}$  and  $\overline{\text{Post}}$  is monotonic, i.e., If  $S_1 \subseteq S_2$  then  $\text{Post}(S_1) \subseteq \text{Post}(S_2)$ .

Exercise. Show that all states that are reachable in exactly  $k$  steps is  $\text{Post}^k(Q_0)$ .

Exercise. If this algorithm terminates and returns  $R$  then  $\text{Reach}_A(Q_0) \subseteq R$ , i.e., it computes an over-approximation of the reachable sets of  $A$ .

$R \cap \text{Unsafe} = \emptyset$  proves safety, but  $\overline{\text{Reach}}_A(Q_0) \cap \text{Unsafe} \neq \emptyset$  does not imply that there is a real counterexample

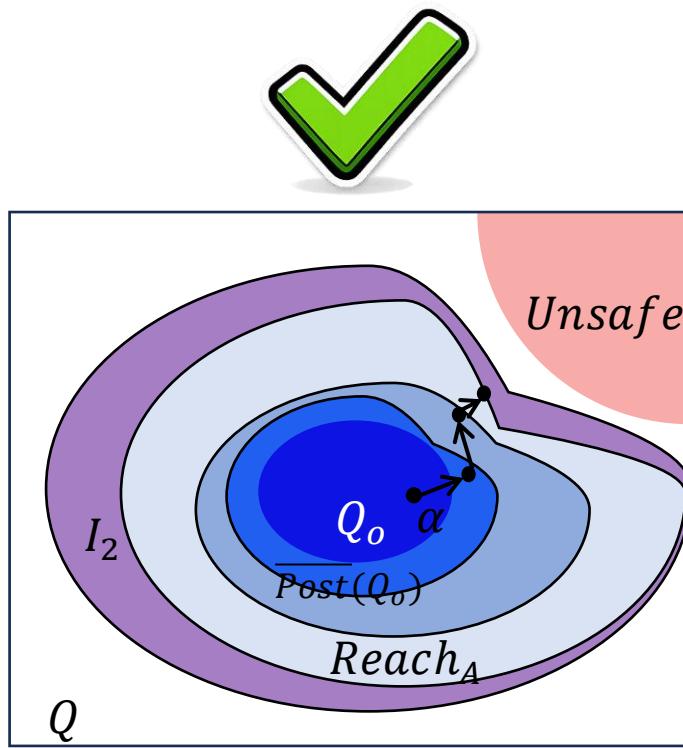


# Invariants and safety verification

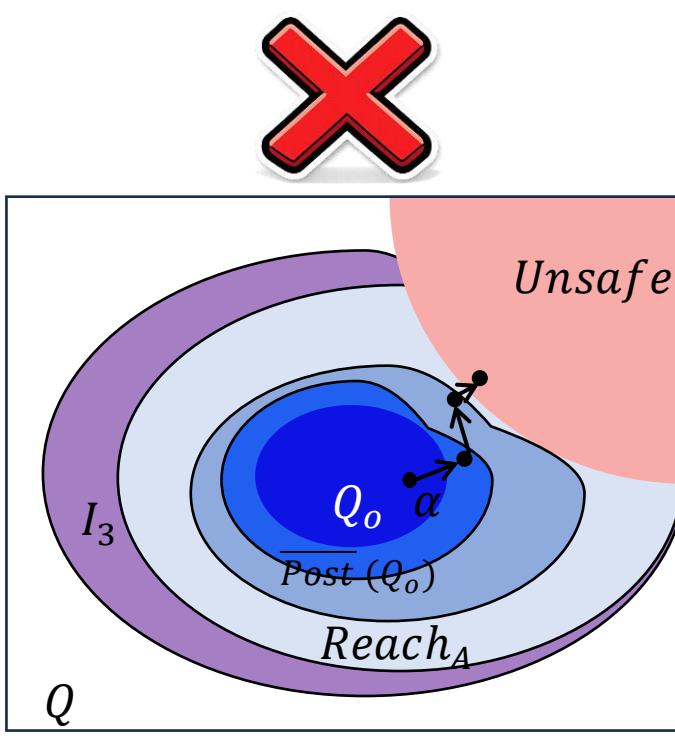
A set  $I \subseteq Q$  is an **invariant** if  $Reach_A(Q_0) \subseteq I$

Over-approximates the reachable states, not unique, and define everything that *can* happen

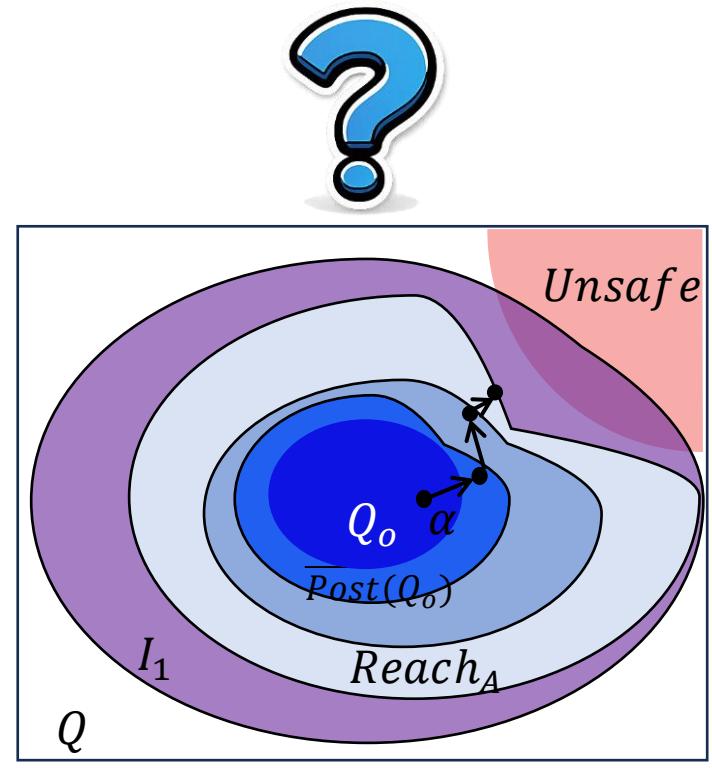
If the algorithm terminates, it returns *an* invariant which may or may not prove safety



System is safe but and verified by invariant  $I_2$



$I_3 \cap Unsafe \neq \emptyset$  and system is unsafe



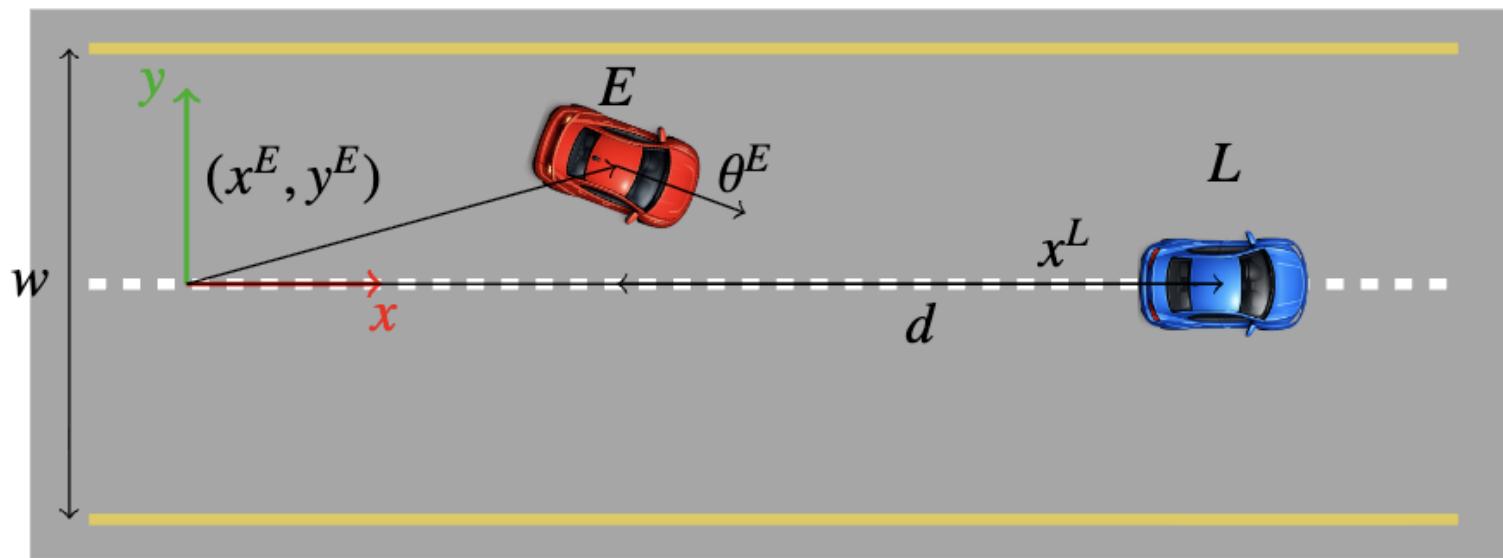
$I_1 \cap Unsafe \neq \emptyset$  but system is safe

# Example: Lane-keeping

Example. Vehicle (E) with braking and lane-keeping controller

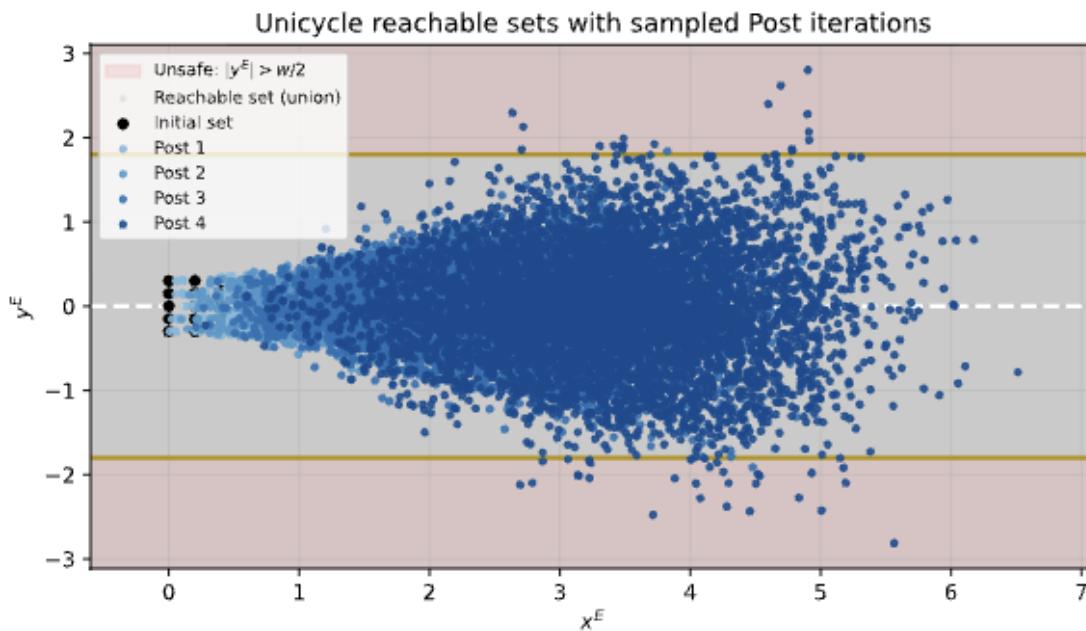
$Q: [x^E, y^E, \theta^E, x^L] \in \mathbb{R}^4$  and velocities are chosen in each step

See course reader for definition of  $D$

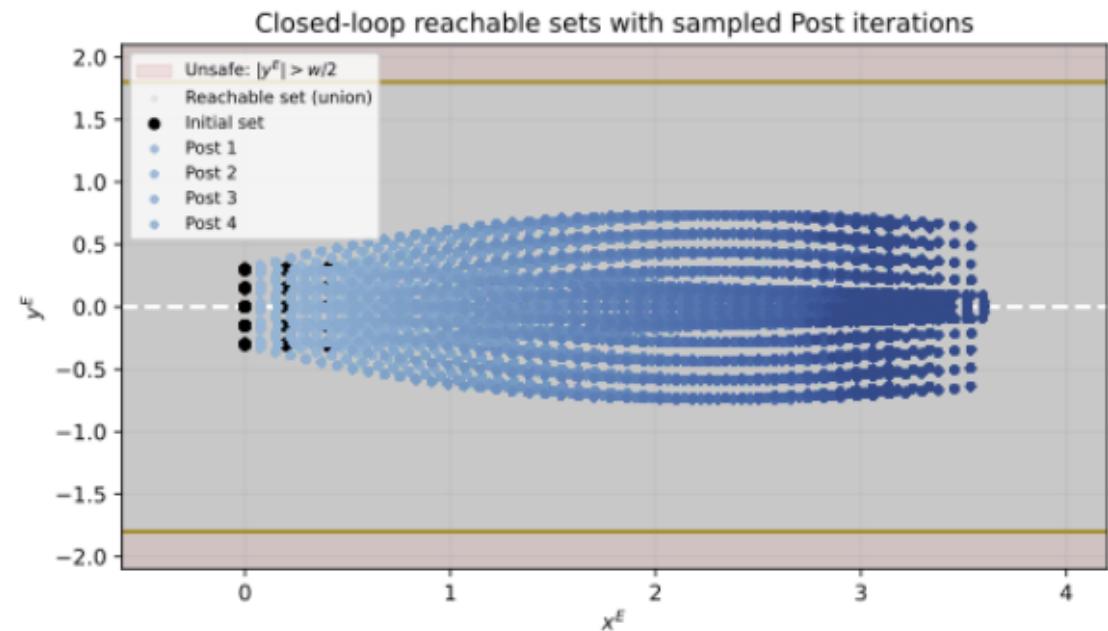


# Reachable sets

## Open-loop



## Closed-loop

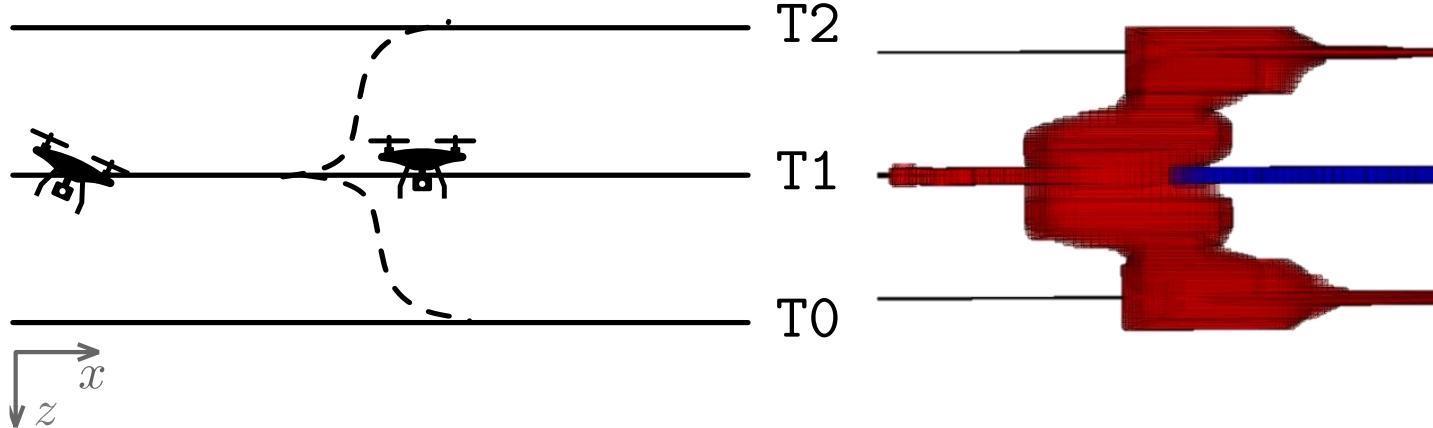


# Summary

- ▶ Canvas quiz:  
[https://canvas.illinois.edu/courses/67113/assignments/1563205?display=full\\_width\\_with\\_nav](https://canvas.illinois.edu/courses/67113/assignments/1563205?display=full_width_with_nav)
- ▶ Verification is the problem of proving/disproving requirements
- ▶ Safety requirements state Unsafe things never happen OR
  - ▶ All reachable states are disjoint from unsafe sets  $\text{Reach}_A \cap \text{Unsafe} = \emptyset$
- ▶ For finite state systems explicit reachability possible via DFS
- ▶ In general, reachability and verification are hard (state space explosion, undecidability)
- ▶ We can over-approximate  $\text{Reach}_A \subseteq \overline{\text{Post}}^k(Q_0)$

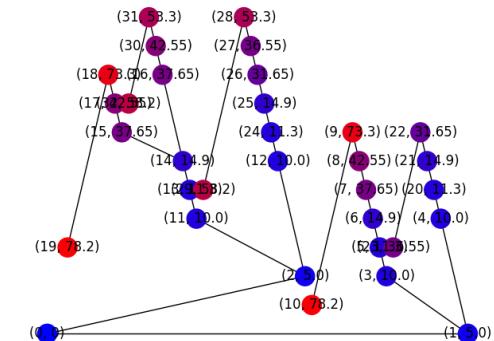
# Verse: Python library for reachability analysis (MPO)

```
class Mode(Enum):
    Normal = auto()
    Up = auto()
...
class Track(Enum):
    T0 = auto()
    T1 = auto()
...
class State:
    x: float
    y: float
...
mode: Mode
track: Track
```



```
def decisionLogic(ego: State, others: List[State], map):
    if ego.mode == Mode.Normal:
        if any(isClose(ego, other) for other in others):
            if map.exist(ego.track, ego.mode, Mode.Up):
                next.mode = Mode.Up
                next.track = map.h(ego.track, ego.mode, Mode.Up)
            if map.exist(ego.track, ego.mode, Mode.Down):
                next.mode = Mode.Down
    ...
assert not any(isVeryClose(ego, other) for other in others), "Separation"
```

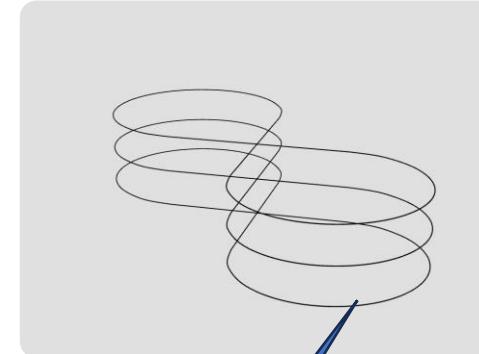
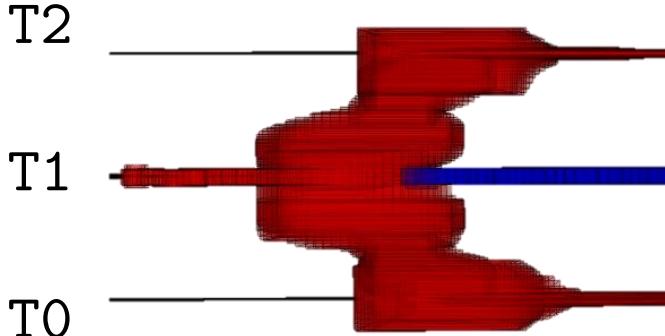
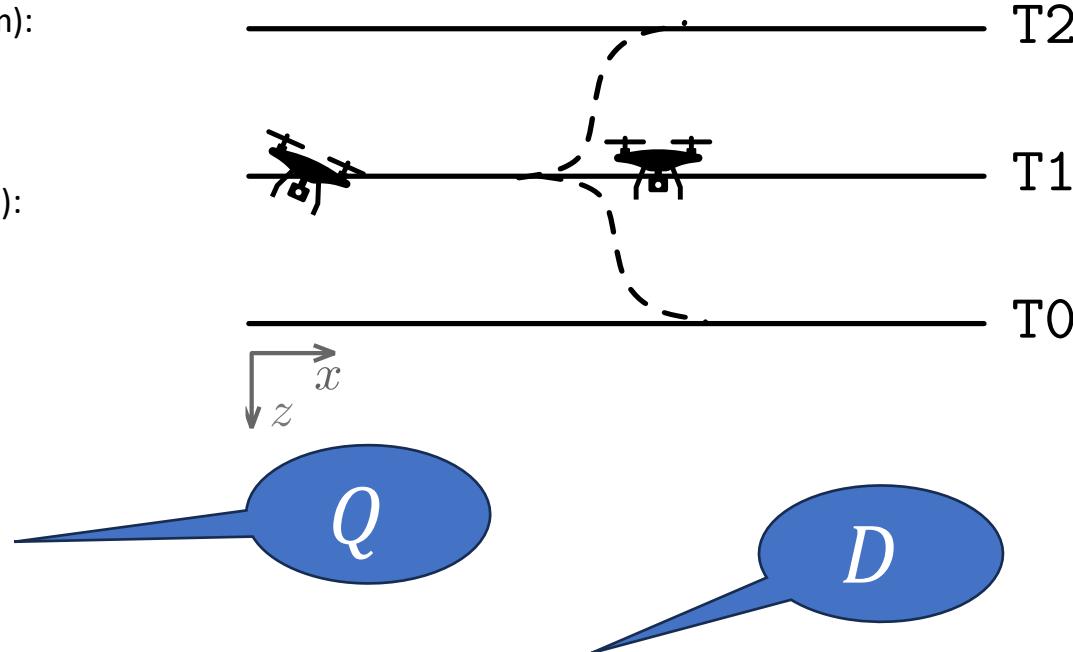
```
q1 = QuadrotorAgent("q1", ...) // Defines the dynamics
q1.set_initial([...], (Mode.Normal, Track.T1))
scenario.add_agent(q1)
q2 = ...
scenario.set_map(M5())
scenario.simulate(...)
scenario.verify(...)
```



# Verse: Python library for reachability analysis (MPO)

```
class Mode(Enum):
    Normal = auto()
    Up = auto()
...
class Track(Enum):
    T0 = auto()
    T1 = auto()
...
class State:
    x: float
    y: float
...
mode: Mode
track: Track
```

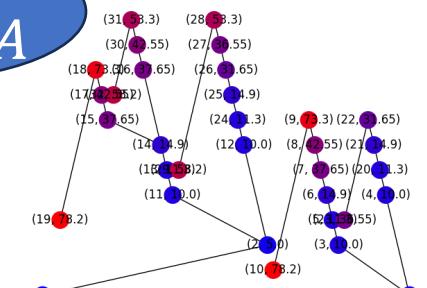
```
def decisionLogic(ego: State, others: List[State], map):
    if ego.mode == Mode.Normal:
        if any(isClose(ego, other) for other in others):
            if map.exist(ego.track, ego.mode, Mode.Up):
                next.mode = Mode.Up
                next.track = map.h(ego.track, ego.mode, Mode.Up)
            if map.exist(ego.track, ego.mode, Mode.Down):
                next.mode = Mode.Down
    ...
assert not any(isVeryClose(ego, other) for other in others), "Separation"
```



```
q1 = QuadrotorAgent("q1", ...)
q1.set_initial(..., (Mode.Normal, Track.T1))
scenario.add_agent(q1)
q2 = ...
scenario.set_map(M5())
scenario.simulate(...)
scenario.verify(...)
```

*Reach<sub>A</sub>*

*Unsafe*



# Inductive invariants

**Proposition 1.** If (i)  $Q_0 \subseteq I$  and (ii)  $\text{Post}(I) \subseteq I$  then  $I$  is an invariant, i.e.,  $\text{Reach}_A \subseteq I$ .

Such invariants are called **inductive invariants**

**Proof.** Consider any reachable state  $\mathbf{q} \in \text{Reach}_A \subseteq Q$

By definition of reachable state, there is an execution  $\alpha$  with  $\alpha_k = \mathbf{q}$

By induction on  $k$  we will show that  $\mathbf{q} \in I$

Base case, for  $k=0$ ,  $\alpha_0 = q_0 \in Q_0 \subseteq I$  [using definition of execution and (i)]

Induction. By inductive hypothesis, suppose  $\alpha_k \in I$ . We have to show  $\mathbf{q} = \alpha_{k+1} \in I$ .

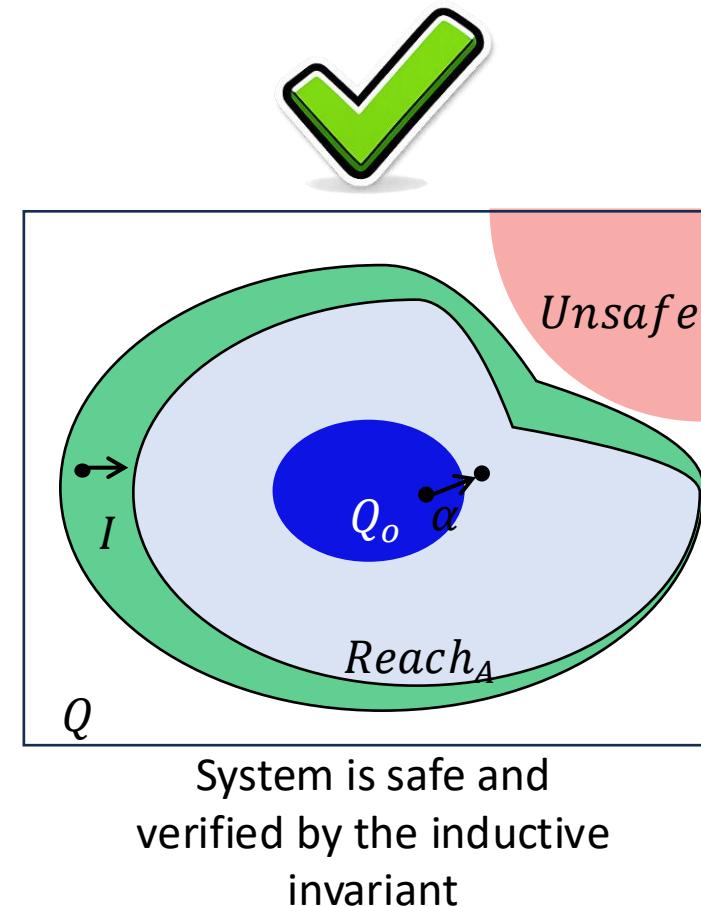
$\mathbf{q} \in \text{Post}(\alpha_k)$  [Definition of Post,  $(\alpha_k, \mathbf{q}) \in D$ ]

$\mathbf{q} \in \text{Post}(I)$  [Monotonicity of Post.  $\alpha_k \in I \Rightarrow \text{Post}(\alpha_k) \subseteq \text{Post}(I)$ ]

$\mathbf{q} \subseteq I$  [By (ii)]

# Inductive invariants and Safety

- ▶ Guess a candidate inductive invariant  $I$
- ▶ If  $I \cap \text{Unsafe} = \emptyset$  and  $Q_0 \subseteq I$  and  $\text{Post}(I) \subseteq I$  then by the Proposition 1  $\text{Reach}_A \subseteq I$  and we have verified safety
- ▶ If the start and transition conditions fail, that does *not* imply that  $I$  is not an invariant
- ▶ It only implies that  $I$  cannot be checked inductively by Proposition 1.



# Revisiting AEB

To prove no crash  $x_2 > x_1$  in all reachable states, we will need assumptions about initial conditions ( $x_{10}, x_{20}, v_{10}, v_{20}$ ), sensing distance ( $d_s$ ), and braking acceleration ( $a_b$ )

Discovering these assumptions (for system correctness) is a valuable side-effect of verification

$$\text{Assumption: } x_{20} - x_{10} > d_s > \frac{v_{10}^2}{ab}$$

The proof of correctness (as expected) will relate total time of braking with the initial separation. We need a timer

# Checking Inductive Invariant for AEB

```

timer = 0
If  $x_2 - x_1 \leq d_s$ 
  If  $v_1 \geq a_b$ 
     $v_1 = v_1 - a_b$ 
    timer := timer+1
  else
     $v_1 = 0$ 
else
   $v_1 = v_1$ 
 $x_2 = x_2 + v_2$ 
 $x_1 = x_1 + v_1$ 

```

Invariant.  $I_1$ :  $\text{timer} + \frac{v_1}{a_b} \leq \frac{v_{10}}{a_b}$ .

Bound on total braking time in terms of velocity and deceleration

Proof. We need to check two conditions for this to be an inductive invariant: (i)  $Q_0 \in I_1$  and (ii)  $\text{Post}(I_1) \subseteq I_1$ .

(i) Consider any  $q \in Q_0$ . We need to show  $q \in I_1$ .

$$q.\text{timer} + \frac{q.v_1}{a_b} = 0 + \frac{v_{10}}{a_b} \leq \frac{v_{10}}{a_b}.$$

(ii) Consider any  $(q, q') \in D$  with  $q \in I_1$ . We need to show  $q' \in I_1$ .

As there are three branches in  $D$ , there are 3 cases.

$$(a) q'.\text{timer} + \frac{q'.v_1}{a_b} = q.\text{timer} + 1 + \frac{q.v_1 - a_b}{a_b} = q.\text{timer} + \frac{q.v_1}{a_b} \leq \frac{v_{10}}{a_b}$$

$$(b) q'.\text{timer} + \frac{q'.v_1}{a_b} = q.\text{timer} + 0 \leq \frac{v_{10}}{a_b}$$

$$(c) q'.\text{timer} + \frac{q'.v_1}{a_b} = q.\text{timer} + \frac{q.v_1}{a_b} \leq \frac{v_{10}}{a_b}$$

$$I_2: \text{timer} \leq \frac{v_{10}}{a_b}$$

# Invariants and assumptions give correctness proof

Consider any two reachable states:

$q_1$  is where  $x_2 - x_1 \leq d_s$  became true first, and

$q_2$  is reached from  $q_1$  with  $q_2 \cdot x_2 - q_2 \cdot x_1 \leq d_s$  (other reachable states are safe)

$$q_2 \cdot x_2 - q_2 \cdot x_1$$

$$> q_1 \cdot x_2 - q_2 \cdot x_1$$

[1, Because  $x_2$  increased]

$$> q_1 \cdot x_2 - q_1 \cdot x_1 - v_{10} \cdot \frac{v_{10}}{a_b}$$

[ $I_2 \Rightarrow \text{timer} \leq \frac{v_{10}}{a_b}$  and  $q_2 \cdot x_1 \leq q_1 \cdot x_1 + v_{10} \cdot \frac{v_{10}}{a_b}$  ]

$$> d_s - \frac{v_{10}^2}{a_b}$$

[By def of  $q_1$  ]

$$> 0$$

[By Assumption]

# Summary

- ▶ Testing alone is inadequate---in theory and practice
- ▶ Automaton (state machine) models, executions, and requirements give us the language to state correctness claims precisely
- ▶ Verification is the problem of proving/disproving such claims
- ▶ Safety claims are a (prevalent) subset of correctness claims
- ▶ Reachability analysis can prove/disprove safety
- ▶ In general, reachability and verification are hard (state space explosion, undecidability)
- ▶ Inductive invariants over-approximating reachable states give a practical method for proving safety